

In [1]: `!pip install yfinance`

```
Requirement already satisfied: yfinance in c:\users\hanshu\anaconda3\lib\site-packages (0.2.65)
Requirement already satisfied: pandas>=1.3.0 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (2.2.2)
Requirement already satisfied: numpy>=1.16.5 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (1.26.4)
Requirement already satisfied: requests>=2.31 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (2.32.3)
Requirement already satisfied: multitasking>=0.0.7 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (0.0.12)
Requirement already satisfied: platformdirs>=2.0.0 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (3.10.0)
Requirement already satisfied: pytz>=2022.5 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (2024.1)
Requirement already satisfied: frozendict>=2.3.4 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (2.4.2)
Requirement already satisfied: peewee>=3.16.2 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (3.18.2)
Requirement already satisfied: beautifulsoup4>=4.11.1 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (4.12.3)
Requirement already satisfied: curl_cffi>=0.7 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (0.13.0)
Requirement already satisfied: protobuf>=3.19.0 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (6.32.0)
Requirement already satisfied: websockets>=13.0 in c:\users\hanshu\anaconda3\lib\site-packages (from yfinance) (15.0.1)
Requirement already satisfied: soupsieve>1.2 in c:\users\hanshu\anaconda3\lib\site-packages (from beautifulsoup4>=4.11.1->yfinance) (2.5)
Requirement already satisfied: cffi>=1.12.0 in c:\users\hanshu\anaconda3\lib\site-packages (from curl_cffi>=0.7->yfinance) (1.17.1)
Requirement already satisfied: certifi>=2024.2.2 in c:\users\hanshu\anaconda3\lib\site-packages (from curl_cffi>=0.7->yfinance) (2025.4.26)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\hanshu\anaconda3\lib\site-packages (from pandas>=1.3.0->yfinance) (2.9.0.post0)
Requirement already satisfied: tzdata>=2022.7 in c:\users\hanshu\anaconda3\lib\site-packages (from pandas>=1.3.0->yfinance) (2023.3)
Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\hanshu\anaconda3\lib\site-packages (from requests>=2.31->yfinance) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in c:\users\hanshu\anaconda3\lib\site-packages (from requests>=2.31->yfinance) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\hanshu\anaconda3\lib\site-packages (from requests>=2.31->yfinance) (2.2.3)
Requirement already satisfied: pycparser in c:\users\hanshu\anaconda3\lib\site-packages (from cffi>=1.12.0->curl_cffi>=0.7->yfinance) (2.21)
Requirement already satisfied: six>=1.5 in c:\users\hanshu\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas>=1.3.0->yfinance) (1.16.0)
```

In [2]: `import yfinance as yf
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor`

```
#from sklearn import metrics
#from sklearn.metrics import accuracy_score, classification_report, confusion_ma
```

In [3]: *#The code fetches historical price data for Bitcoin, Ethereum, Tether, and Binan*  
*#This cleaned data can then be used for further analysis or machine Learning tas*

```
btc = yf.Ticker('BTC-USD')
prices1 = btc.history(period='5y')
prices1.drop(columns=['Open', 'High', 'Low', 'Dividends', 'Stock Splits'], axis=

eth = yf.Ticker('ETH-USD')
prices2 =eth.history(period='5y')
prices2.drop(columns=['Open', 'High', 'Low', 'Dividends', 'Stock Splits'], axis

usdt = yf.Ticker('USDT-USD')
prices3 = usdt.history(period='5y')
prices3.drop(columns=['Open', 'High', 'Low', 'Dividends', 'Stock Splits'], axis

bnb = yf.Ticker('BNB-USD')
prices4 = bnb.history(period='5y')
prices4.drop(columns=['Open', 'High', 'Low', 'Dividends', 'Stock Splits'], axis
```

In [4]: *#The parameters lsuffix and rsuffix in the join method are used to add suffixes*  
*# This is necessary to avoid column name conflicts when the two DataFrames have*


```
p1 = prices1.join(prices2, lsuffix= '(BTC)', rsuffix= '(ETH)')
p2 = prices3.join(prices4, lsuffix = '(USDT)', rsuffix = '(BNB)')
data = p1.join(p2, lsuffix = '_', rsuffix = '_')
```

In [5]: data.head()

Out[5]:

	Close(BTC)	Volume(BTC)	Close(ETH)	Volume(ETH)	Close(USDT)	Vo
--	------------	-------------	------------	-------------	-------------	----

Date						
2020-09-14 00:00:00+00:00	10680.837891	35453581940	377.268860	17536695361	1.001289	4
2020-09-15 00:00:00+00:00	10796.951172	32509451925	364.839203	16140584321	1.002487	4
2020-09-16 00:00:00+00:00	10974.905273	30769986455	365.812286	16107612177	1.003444	5
2020-09-17 00:00:00+00:00	10948.990234	38151810523	389.019226	19899531080	1.001878	5
2020-09-18 00:00:00+00:00	10944.585938	26341903912	384.364532	14108357740	0.999502	4



In [6]: data.tail()

Out[6]:

	Close(BTC)	Volume(BTC)	Close(ETH)	Volume(ETH)	Close(USDT)	
Date						
2025-09-10 00:00:00+00:00	113955.359375	56377473784	4349.145996	39521365146	1.000138	
2025-09-11 00:00:00+00:00	115507.539062	45685065332	4461.233398	35959212991	1.000266	
2025-09-12 00:00:00+00:00	116101.578125	54785725894	4715.246094	43839753626	1.000618	
2025-09-13 00:00:00+00:00	115950.507812	34549454947	4668.179688	34843845977	1.000319	
2025-09-14 00:00:00+00:00	115852.859375	31612506112	4642.932617	28678072320	1.000401	

In [7]:

data.shape

Out[7]: (1827, 8)

In [8]:

data.info()  
  
<class 'pandas.core.frame.DataFrame'>  
DatetimeIndex: 1827 entries, 2020-09-14 00:00:00+00:00 to 2025-09-14 00:00:00+00:00  
Data columns (total 8 columns):  
# Column Non-Null Count Dtype  
--- ---  
0 Close(BTC) 1827 non-null float64  
1 Volume(BTC) 1827 non-null int64  
2 Close(ETH) 1827 non-null float64  
3 Volume(ETH) 1827 non-null int64  
4 Close(USDT) 1827 non-null float64  
5 Volume(USDT) 1827 non-null int64  
6 Close(BNB) 1827 non-null float64  
7 Volume(BNB) 1827 non-null int64  
dtypes: float64(4), int64(4)  
memory usage: 128.5 KB

In [9]:

data.isna().sum()

Out[9]: Close(BTC) 0  
Volume(BTC) 0  
Close(ETH) 0  
Volume(ETH) 0  
Close(USDT) 0  
Volume(USDT) 0  
Close(BNB) 0  
Volume(BNB) 0  
dtype: int64

In [10]:

data.describe()

Out[10]:

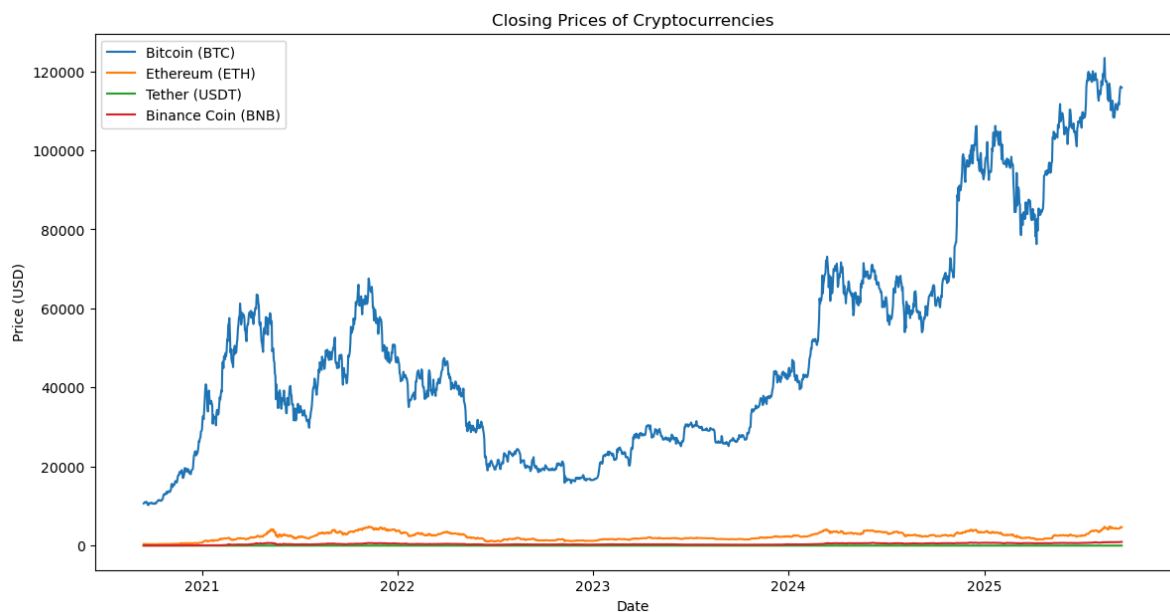
	Close(BTC)	Volume(BTC)	Close(ETH)	Volume(ETH)	Close(USDT)	Volume(U
<b>count</b>	1827.000000	1.827000e+03	1827.000000	1.827000e+03	1827.000000	1.827000
<b>mean</b>	49323.487513	3.571168e+10	2350.643495	1.842279e+10	1.000188	6.659633
<b>std</b>	28103.797749	2.156630e+10	983.483850	1.191476e+10	0.000737	4.142044
<b>min</b>	10246.186523	5.331173e+09	321.116302	2.081626e+09	0.995872	9.989859
<b>25%</b>	27055.889648	2.136936e+10	1649.178711	1.024022e+10	0.999922	3.915764
<b>50%</b>	42412.433594	3.115874e+10	2260.648682	1.581725e+10	1.000157	5.707433
<b>75%</b>	63842.345703	4.430770e+10	3107.366699	2.290538e+10	1.000430	8.167255
<b>max</b>	123344.062500	3.509679e+11	4831.348633	9.245355e+10	1.011530	3.006686



### Exploratory Data Analysis

In [16]: *#Visualize the Closing Prices*  
*# create a line plot to visualize the closing prices of all four cryptocurrencie*

```
plt.figure(figsize=(14,7))
plt.plot(data.index, data['Close(BTC)'], label='Bitcoin (BTC)')
plt.plot(data.index, data['Close(ETH)'], label='Ethereum (ETH)')
plt.plot(data.index, data['Close(USDT)'], label='Tether (USDT)')
plt.plot(data.index, data['Close(BNB)'], label='Binance Coin (BNB)')
plt.title('Closing Prices of Cryptocurrencies')
plt.xlabel('Date')
plt.ylabel('Price (USD)')
plt.legend()
plt.show()
```




In [15]: `print(data.columns)`

```
Index(['Close(BTC)', 'Volume(BTC)', 'Close(ETH)', 'Volume(ETH)', 'Close(USDT)',
      'Volume(USDT)', 'Close(BNB)', 'Volume(BNB)'],
      dtype='object')
```

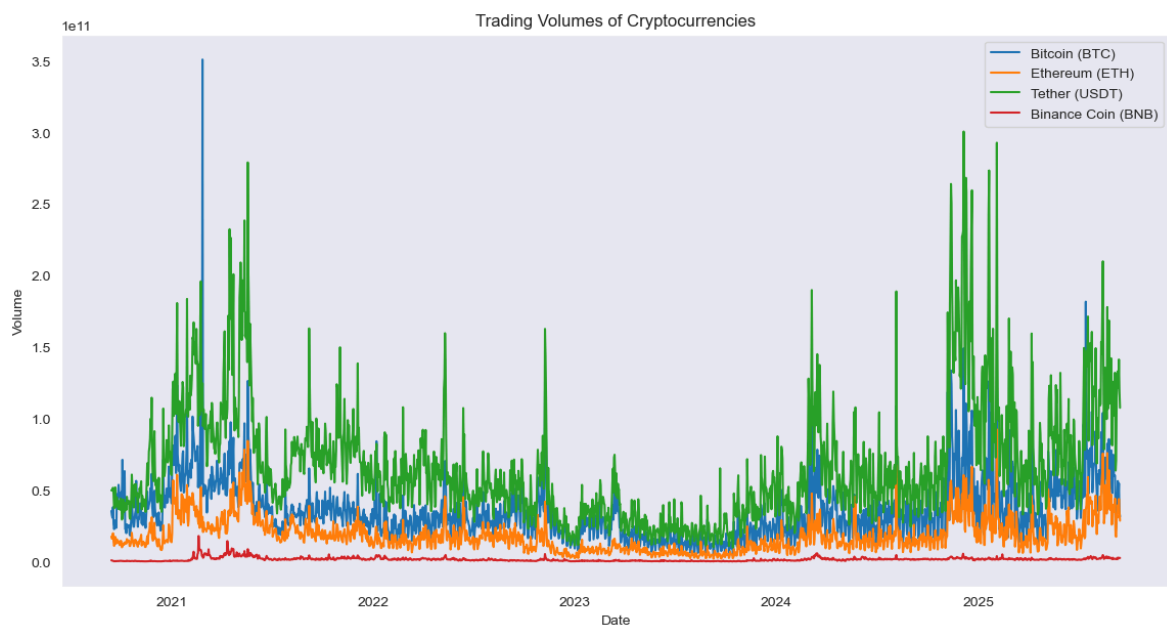
```
In [18]: plt.figure(figsize = (25,5))
sns.set_style('dark')
sns.lineplot(data=data)
```

Out[18]: <Axes: xlabel='Date'>



```
In [21]: # Visualize the Trading Volumes
#Let's visualize the trading volumes of all four cryptocurrencies:
```

```
plt.figure(figsize=(14, 7))
plt.plot(data.index, data['Volume(BTC)'], label='Bitcoin (BTC)')
plt.plot(data.index, data['Volume(ETH)'], label='Ethereum (ETH)')
plt.plot(data.index, data['Volume(USDT)'], label='Tether (USDT)')
plt.plot(data.index, data['Volume(BNB)'], label='Binance Coin (BNB)')
plt.title('Trading Volumes of Cryptocurrencies')
plt.xlabel('Date')
plt.ylabel('Volume')
plt.legend()
plt.show()
```



```
In [22]: #Correlation Analysis
#We'll analyze the correlation between the closing prices of the cryptocurrencies
# Calculate the correlation matrix
```

```
corr_matrix = data[['Close(BTC)', 'Close(ETH)', 'Close(USDT)', 'Close(BNB)']].corr()

# plot the heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', vmin=1, vmax=1)
plt.title('Correlation Matrix of Closing Prices')
plt.show()
```



```
In [24]: # Distribution of Closing Prices
#Let's plot the distribution of closing prices for each cryptocurrency:

plt.figure(figsize=(14,7))

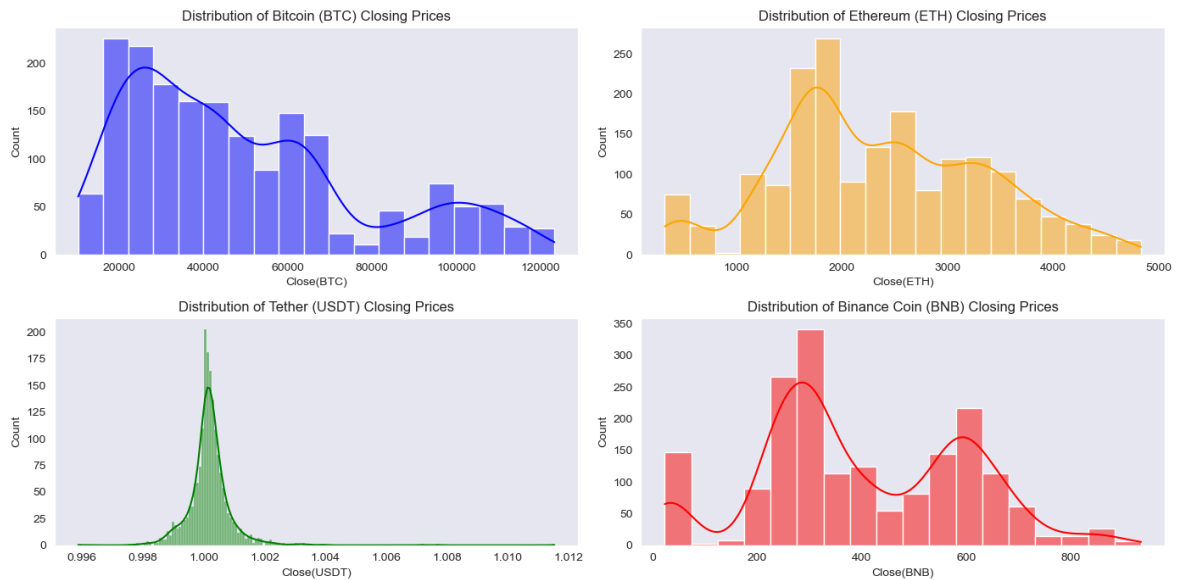
plt.subplot(2, 2, 1)
sns.histplot(data['Close(BTC)'], kde=True, color='blue')
plt.title('Distribution of Bitcoin (BTC) Closing Prices')

plt.subplot(2, 2, 2)
sns.histplot(data['Close(ETH)'], kde=True, color='orange')
plt.title('Distribution of Ethereum (ETH) Closing Prices')

plt.subplot(2, 2, 3)
sns.histplot(data['Close(USDT)'], kde=True, color='green')
plt.title('Distribution of Tether (USDT) Closing Prices')

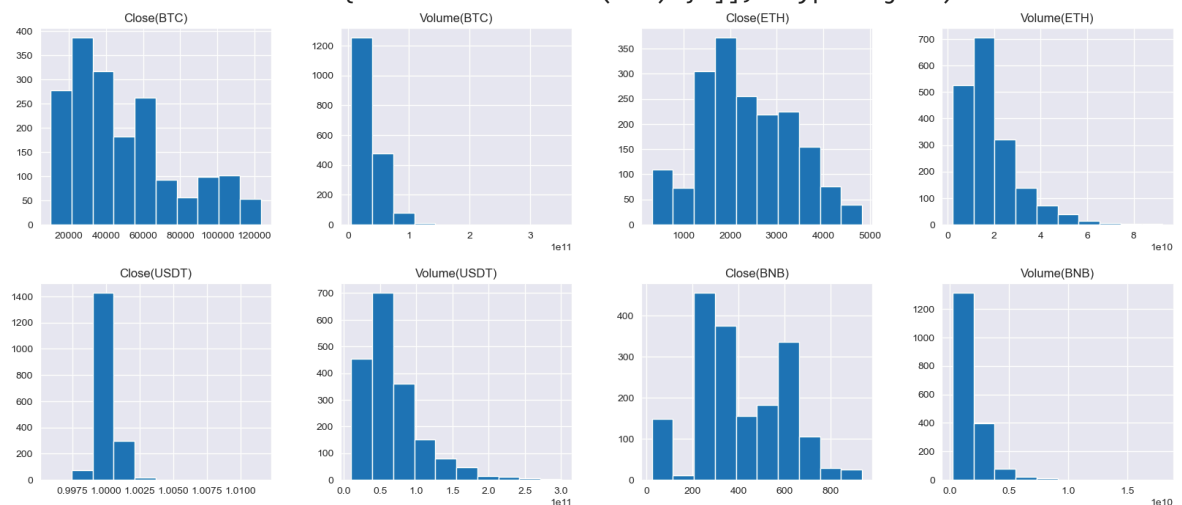
plt.subplot(2, 2, 4)
sns.histplot(data['Close(BNB)'], kde=True, color='red')
plt.title('Distribution of Binance Coin (BNB) Closing Prices')

plt.tight_layout()
plt.show()
```



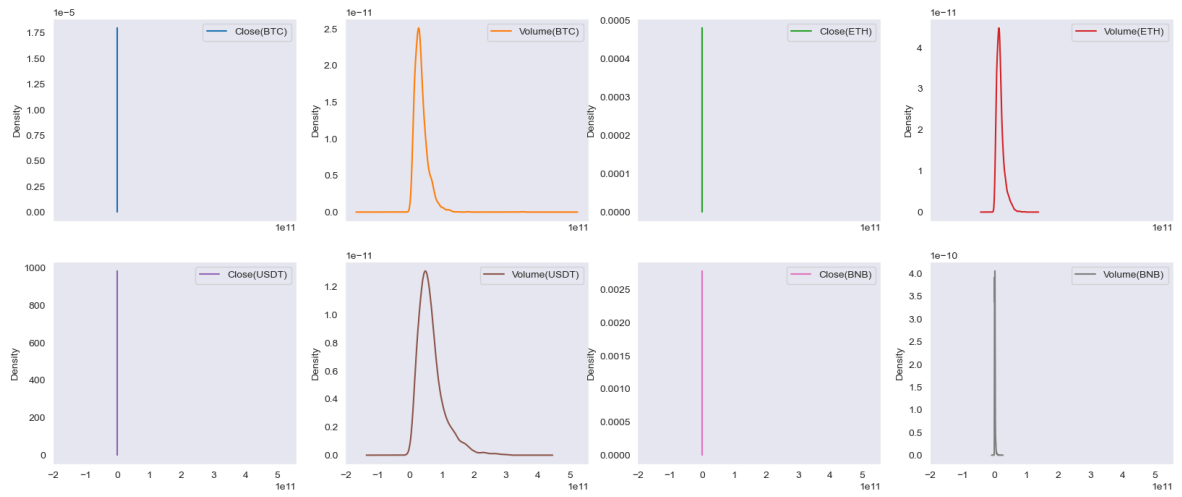
```
In [25]: data.hist(figsize=(20,8), layout=(2,4))
```

```
Out[25]: array([[<Axes: title={'center': 'Close(BTC)'}>,
  <Axes: title={'center': 'Volume(BTC)'}>,
  <Axes: title={'center': 'Close(ETH)'}>,
  <Axes: title={'center': 'Volume(ETH)'}>],
  [<Axes: title={'center': 'Close(USDT)'}>,
  <Axes: title={'center': 'Volume(USDT)'}>,
  <Axes: title={'center': 'Close(BNB)'}>,
  <Axes: title={'center': 'Volume(BNB)'}>]], dtype=object)
```



```
In [26]: data.plot(kind = "kde", subplots = True, layout = (2, 4), figsize = (20, 8))
```

```
Out[26]: array([[<Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
  <Axes: ylabel='Density'>, <Axes: ylabel='Density'>],
  [<Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
  <Axes: ylabel='Density'>, <Axes: ylabel='Density'>]], dtype=object)
```



```
In [28]: sns.pairplot(data.sample(n=100));
```



```
In [29]: X = data.drop(columns = ['Close(BTC)'], axis = 1)
Y = data.loc[:, 'Close(BTC)']
```

```
In [30]: X.head()
```



Out[30]:

	Volume(BTC)	Close(ETH)	Volume(ETH)	Close(USDT)	Volume(USDT)	C
Date						
2020-09-14 00:00:00+00:00	35453581940	377.268860	17536695361	1.001289	49936255991	
2020-09-15 00:00:00+00:00	32509451925	364.839203	16140584321	1.002487	49718173930	
2020-09-16 00:00:00+00:00	30769986455	365.812286	16107612177	1.003444	50682289026	
2020-09-17 00:00:00+00:00	38151810523	389.019226	19899531080	1.001878	51695424541	
2020-09-18 00:00:00+00:00	26341903912	384.364532	14108357740	0.999502	47248825663	

In [31]: X.tail()

Out[31]:

	Volume(BTC)	Close(ETH)	Volume(ETH)	Close(USDT)	Volume(USDT)
Date					
2025-09-10 00:00:00+00:00	56377473784	4349.145996	39521365146	1.000138	133101421364
2025-09-11 00:00:00+00:00	45685065332	4461.233398	35959212991	1.000266	121507255807
2025-09-12 00:00:00+00:00	54785725894	4715.246094	43839753626	1.000618	141338448172
2025-09-13 00:00:00+00:00	34549454947	4668.179688	34843845977	1.000319	119042646333
2025-09-14 00:00:00+00:00	31612506112	4642.932617	28678072320	1.000401	107577327616

In [32]: Y.head()

Out[32]:

```
Date
2020-09-14 00:00:00+00:00    10680.837891
2020-09-15 00:00:00+00:00    10796.951172
2020-09-16 00:00:00+00:00    10974.905273
2020-09-17 00:00:00+00:00    10948.990234
2020-09-18 00:00:00+00:00    10944.585938
Name: Close(BTC), dtype: float64
```

In [33]:

```
# Split the data into training and testing sets

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_
```

In [34]:

```
# Print the shapes of the resulting datasets

print(f'x_train shape: {x_train.shape}')
print(f'x_test shape: {x_test.shape}')
```

```
print(f'y_train shape: {y_train.shape}')
print(f'y_test shape: {y_test.shape}')
```

```
x_train shape: (1461, 7)
x_test shape: (366, 7)
y_train shape: (1461,)
y_test shape: (366,)
```

```
In [35]: #SelectKBest
#SelectKBest is a feature selection method provided by scikit-learn (sklearn) th
#This function evaluates each feature independently and selects those that have

#Parameters
#k: Specifies the number of top features to select. In your case, k=4 indicates

from sklearn.feature_selection import SelectKBest

fs = SelectKBest(k=4)
x_train = fs.fit_transform(x_train, y_train)
x_test = fs.transform(x_test)
```

```
c:\Users\Hanshu\anaconda3\Lib\site-packages\sklearn\feature_selection\_univariate
_selection.py:109: RuntimeWarning: invalid value encountered in divide
msw = ssw / float(dfwn)
```

```
In [37]: mask = fs.get_support()
selected_features = X.columns[mask]
print('Selected Features:', selected_features)
```

```
Selected Features: Index(['Close(USDT)', 'Volume(USDT)', 'Close(BNB)', 'Volume(BN
B)'], dtype='object')
```

```
In [38]: x_train
```

```
Out[38]: array([[1.00051105e+00, 6.22330509e+10, 3.45933685e+02, 2.26992368e+09],
 [9.99966025e-01, 9.37161581e+10, 5.99706543e+02, 6.58591900e+09],
 [9.99065995e-01, 4.09658641e+10, 3.19609009e+02, 1.60429741e+09],
 ...,
 [1.00018799e+00, 5.84567573e+10, 6.29942871e+02, 2.23811852e+09],
 [1.00037599e+00, 5.91036805e+10, 4.30503265e+02, 1.49950509e+09],
 [1.00028896e+00, 5.74898269e+10, 2.87536133e+02, 1.57890533e+09]])
```

```
In [39]: #MinMaxScaler is a preprocessing method in scikit-learn that transforms features
# It's often used when your data needs to be normalized within a specific range

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

```
In [42]: # implementation of 10 different regression algorithms using scikit-learn. Each

#Import Libraries and Generate Sample Data

from sklearn.datasets import make_regression
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
```

```
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [49]: #Define Models and Perform Training and Evaluation

models = {
    'Linear Regression': LinearRegression(),
    'Ridge Regression': Ridge(alpha=1.0),
    'Lasso Regression': Lasso(alpha=1.0),
    'ElasticNet Regression': ElasticNet(alpha=1.0, l1_ratio=0.5),
    'Support Vector Regression (SVR)': SVR(kernel='rbf'),
    'Decision Tree Regression': DecisionTreeRegressor(),
    'Random Forest Regression': RandomForestRegressor(n_estimators=100),
    'Gradient Boosting Regression': GradientBoostingRegressor(n_estimators=100),
    'K-Neighbors Regression': KNeighborsRegressor(n_neighbors=5),
    'Neural Network (MLP) Regression': MLPRegressor(hidden_layer_sizes=(100, 50)
}

# Train and evaluate each model
results = {'Model': [], 'MSE': [], 'R-squared': []}

for name, model in models.items():
    # Train the model
    model.fit(x_train, y_train)

    # Predict on test set
    y_pred = model.predict(x_test)

    # Evaluate model
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    # Store resultsre
    results['Model'].append(name)
    results['MSE'].append(mse)
    results['R-squared'].append(r2)

    # print results
    print(f"----- {name} -----")
    print(f"Mean Squared Error (MSE): {mse}")
    print(f"R-squared: {r2}")
    print()

# Convert results to DataFrame for visualization
results_df = pd.DataFrame(results)
print(results_df)

# Plotting the results
plt.figure(figsize=(12, 6))
plt.barh(results_df['Model'], results_df['R-squared'], color='skyblue')
plt.xlabel('R-squared')
plt.title('R-squared of Different Regression Models')
plt.xlim(-1, 1)
plt.gca().invert_yaxis()
plt.show()
```

----- Linear Regression -----  
Mean Squared Error (MSE): 150332109.8236788  
R-squared: 0.8086771935356407

----- Ridge Regression -----  
Mean Squared Error (MSE): 151893302.8553303  
R-squared: 0.8066903137359858

----- Lasso Regression -----  
Mean Squared Error (MSE): 150357484.84905022  
R-squared: 0.8086448995628237

----- ElasticNet Regression -----  
Mean Squared Error (MSE): 671371106.5323241  
R-squared: 0.14556774044147236

----- Support Vector Regression (SVR) -----  
Mean Squared Error (MSE): 813376297.152881  
R-squared: -0.035157665686978756

----- Decision Tree Regression -----  
Mean Squared Error (MSE): 85266293.85579413  
R-squared: 0.8914843498409029

----- Random Forest Regression -----  
Mean Squared Error (MSE): 39904889.378470555  
R-squared: 0.9492143399271561


----- Gradient Boosting Regression -----  
Mean Squared Error (MSE): 50841728.860370405  
R-squared: 0.9352953785955983

----- K-Neighbors Regression -----  
Mean Squared Error (MSE): 49073045.421996176  
R-squared: 0.9375463247146522

c:\Users\Hanshu\anaconda3\Lib\site-packages\sklearn\normalization\\_multilayer\_perceptron.py:690: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.  
warnings.warn(

----- Neural Network (MLP) Regression -----  
Mean Squared Error (MSE): 495536206.3573425  
R-squared: 0.36934712207401244

	Model	MSE	R-squared
0	Linear Regression	1.503321e+08	0.808677
1	Ridge Regression	1.518933e+08	0.806690
2	Lasso Regression	1.503575e+08	0.808645
3	ElasticNet Regression	6.713711e+08	0.145568
4	Support Vector Regression (SVR)	8.133763e+08	-0.035158
5	Decision Tree Regression	8.526629e+07	0.891484
6	Random Forest Regression	3.990489e+07	0.949214
7	Gradient Boosting Regression	5.084173e+07	0.935295
8	K-Neighbors Regression	4.907305e+07	0.937546
9	Neural Network (MLP) Regression	4.955362e+08	0.369347



In [ ]:

```
In [52]: import pickle
import numpy as np
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, r2_score

# Generate sample data
X, Y = make_regression(n_samples=1000, n_features=10, noise=0.1, random_state=0)

# Scale the features (optional but recommended for some algorithms)
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Initialize Random Forest Regressor
model_rf = RandomForestRegressor(n_estimators=100, random_state=0)

# Train the model
model_rf.fit(X_train, y_train)

# Save the model to a file
filename = 'random_forest_model.pkl'
pickle.dump(model_rf, open(filename, 'wb'))

# Save scaler to a file
with open('scaler.pkl', 'wb') as f:
    pickle.dump(scaler, f)

# Load the model from the file
loaded_model = pickle.load(open(filename, 'rb'))

# Predict using the loaded model
y_pred = loaded_model.predict(X_test)

# Evaluate the loaded model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
print(f"Loaded Random Forest Regression - Mean Squared Error (MSE): {mse}")  
print(f"Loaded Random Forest Regression - R-squared: {r2}")
```

```
Loaded Random Forest Regression - Mean Squared Error (MSE): 38784117.98190298  
Loaded Random Forest Regression - R-squared: 0.9506407093784183
```

In [ ]: